



HealthFlix Machine Learning Project

Project By : PRASAD JADHAV

```
In [1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

```
In [12]: # Load Dataset
file_path = 'healthcare_dataset.csv'
df = pd.read_csv(file_path)
pd.set_option('display.max_columns',30)
print(df.shape)
```

(55500, 15)

```
In [22]: df.head()
```

Out[22]:

	Name	Age	Gender	Blood Type	Medical Condition	Date of Admission	Doctor	Hospital	Insurance Provider	Billing Amount
0	Bobby Jackson	30	Male	B-	Cancer	2024-01-31	Matthew Smith	Sons and Miller	Blue Cross	18856.281301
1	Leslie Terry	62	Male	A+	Obesity	2019-08-20	Samantha Davies	Kim Inc	Medicare	33643.32728
2	Danny Smith	76	Female	A-	Obesity	2022-09-22	Tiffany Mitchell	Cook PLC	Aetna	27955.09607
3	Andrew Watts	28	Female	O+	Diabetes	2020-11-18	Kevin Wells	Hernandez Rogers and Vang,	Medicare	37909.782411
4	Adrienne Bell	43	Female	AB+	Cancer	2022-09-19	Kathleen Hanna	White-White	Aetna	14238.31781

In [5]: `df.tail()`

Out[5]:

	Name	Age	Gender	Blood Type	Medical Condition	Date of Admission	Doctor	Hospital	Insurance Provider
55495	Elizabeth Jackson	42	Female	O+	Asthma	2020-08-16	Joshua Jarvis	Jones-Thompson	Blue Cross
55496	Kyle Perez	61	Female	AB-	Obesity	2020-01-23	Taylor Sullivan	Tucker-Moyer	Cigna
55497	Heather Wang	38	Female	B+	Hypertension	2020-07-13	Joe Jacobs DVM	and Mahoney Johnson Vasquez,	UnitedHealthcare
55498	Jennifer Jones	43	Male	O-	Arthritis	2019-05-25	Kimberly Curry	Jackson Todd and Castro,	Medicare
55499	James Garcia	53	Female	O+	Arthritis	2024-04-02	Dennis Warren	Henry Sons and	Aetna

Data Exploration

In [6]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55500 entries, 0 to 55499
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Name                   55500 non-null  object
1   Age                    55500 non-null  int64
2   Gender                 55500 non-null  object
3   Blood Type             55500 non-null  object
4   Medical Condition      55500 non-null  object
5   Date of Admission      55500 non-null  object
6   Doctor                 55500 non-null  object
7   Hospital               55500 non-null  object
8   Insurance Provider     55500 non-null  object
9   Billing Amount          55500 non-null  float64
10  Room Number            55500 non-null  int64
11  Admission Type         55500 non-null  object
12  Discharge Date         55500 non-null  object
13  Medication              55500 non-null  object
14  Test Results           55500 non-null  object
dtypes: float64(1), int64(2), object(12)
memory usage: 6.4+ MB

```

```
In [7]: df.isnull().sum()
```

```

Out[7]: Name                0
        Age                0
        Gender             0
        Blood Type         0
        Medical Condition   0
        Date of Admission   0
        Doctor              0
        Hospital            0
        Insurance Provider   0
        Billing Amount       0
        Room Number         0
        Admission Type      0
        Discharge Date      0
        Medication          0
        Test Results        0
        dtype: int64

```

```
In [13]: df.duplicated().sum()
```

```
Out[13]: 534
```

```
In [5]: # df = df.drop_duplicates()
```

```
In [12]: df.describe()
```

Out[12]:

	Age	Billing Amount	Room Number
count	54966.000000	54966.000000	54966.000000
mean	51.535185	25544.306284	301.124404
std	19.605661	14208.409711	115.223143
min	13.000000	-2008.492140	101.000000
25%	35.000000	13243.718641	202.000000
50%	52.000000	25542.749145	302.000000
75%	68.000000	37819.858159	401.000000
max	89.000000	52764.276736	500.000000

```
In [13]: cat_cols = [x for x in df.columns if df[x].dtypes != 'float64']

for col in cat_cols:
    print(f"Value counts for column '{col}':")
    print(df[col].value_counts())
    print("\n" + "_"*40 + "\n")
```

```
Value counts for column 'Name':
Name
DAvId muNoZ          3
kaTheRIne WeBSTer    2
mICHael aNdERSon     2
DaVID caLhouN        2
MELiSsA COloN        2
..
dUstin blaCKwELl     1
MARc CLaRK           1
sTEphen AyaLa        1
ThOMaS torreS        1
HARoLD ACOSTa        1
Name: count, Length: 49992, dtype: int64
```

```
Value counts for column 'Age':
Age
38    890
57    881
37    880
34    858
80    855
...
88     25
16     24
14     18
13     14
89      8
Name: count, Length: 77, dtype: int64
```

```
Value counts for column 'Gender':
Gender
Male    27496
Female  27470
Name: count, dtype: int64
```

```
Value counts for column 'Blood Type':
Blood Type
A-    6898
A+    6896
B+    6885
AB+   6882
AB-   6874
B-    6872
O+    6855
O-    6804
Name: count, dtype: int64
```

Value counts for column 'Medical Condition':

Medical Condition

Arthritis 9218

Diabetes 9216

Hypertension 9151

Obesity 9146

Cancer 9140

Asthma 9095

Name: count, dtype: int64

Value counts for column 'Date of Admission':

Date of Admission

2024-03-16 50

2020-10-22 49

2021-12-28 48

2023-08-10 47

2022-07-24 47

..

2022-05-28 14

2023-04-12 14

2022-05-23 13

2019-07-22 13

2022-02-05 12

Name: count, Length: 1827, dtype: int64

Value counts for column 'Doctor':

Doctor

Michael Smith 27

John Smith 22

Robert Smith 21

James Smith 20

Michael Johnson 20

..

Shane Tate 1

Christy Parker 1

Larry Miller 1

Chelsea Neal 1

Jeffrey Moore 1

Name: count, Length: 40341, dtype: int64

Value counts for column 'Hospital':

Hospital

LLC Smith 44

Ltd Smith 39

Johnson PLC 37

Smith Ltd 37

Smith Group 36

..

PLC Navarro 1

PLC McIntosh 1

```
and Hernandez, Hughes Walton      1
Myers-Williams                    1
Moreno Murphy, Griffith and      1
Name: count, Length: 39876, dtype: int64
```

Value counts for column 'Insurance Provider':

```
Insurance Provider
Cigna                11139
Medicare             11039
UnitedHealthcare    11014
Blue Cross           10952
Aetna                10822
Name: count, dtype: int64
```

Value counts for column 'Room Number':

```
Room Number
393      176
104      174
420      174
491      173
209      170
...
189      112
257      111
381      110
254      108
398      108
Name: count, Length: 400, dtype: int64
```

Value counts for column 'Admission Type':

```
Admission Type
Elective      18473
Urgent        18391
Emergency     18102
Name: count, dtype: int64
```

Value counts for column 'Discharge Date':

```
Discharge Date
2020-03-15     53
2021-12-13     51
2023-04-29     51
2020-12-02     50
2020-08-11     50
..
2024-06-04      2
2024-06-05      2
2019-05-11      2
2019-05-09      1
```

2024-06-06 1
Name: count, Length: 1856, dtype: int64

Value counts for column 'Medication':

Medication	
Lipitor	11038
Ibuprofen	11023
Aspirin	10984
Paracetamol	10965
Penicillin	10956

Name: count, dtype: int64

Value counts for column 'Test Results':

Test Results	
Abnormal	18437
Normal	18331
Inconclusive	18198

Name: count, dtype: int64

```
In [14]: num_cols = [x for x in df.columns if df[x].dtypes == 'float64']

for col in num_cols:
    print(f"Value counts for column '{col}':")
    print(df[col].value_counts())
    print("\n" + "_"*40 + "\n")
```

Value counts for column 'Billing Amount':

Billing Amount	
8926.285937	2
8693.755844	2
17889.765079	2
30679.871088	2
1709.059684	2
..	
46506.415756	1
5343.806298	1
17180.108948	1
47078.702712	1
40116.177618	1

Name: count, Length: 50000, dtype: int64


```
In [15]: num_features = df.select_dtypes(include = ['int64', 'float64']).dtypes.index
```

```
In [16]: plt.figure(figsize=(15,15))
plt.suptitle('Univariate Analysis of Features',fontweight='bold',fontsize=15)
```



```
for i in range(0,len(num_features)):
    plt.subplot(10,4,i+1)
    sns.kdeplot(x=df[num_features[i]],shade=True,color='#03045E')
    plt.tight_layout()
```

Univariate Analysis of Features



```
In [17]: plt.figure(figsize = (15,15))
plt.suptitle('Univariate Analysis of Features',fontweight='bold',fontsize=20)

for i in range(0,len(num_features)):
    plt.subplot(10,5,i+1)
    sns.boxplot(data=df,x=num_features[i],color='#008000')
    plt.xlabel(num_features[i])
    plt.tight_layout()
```

Univariate Analysis of Features



Feature Engineering & Preprocessing

```
In [6]: # Drop irrelevant columns
df = df.drop(columns=['Name', 'Doctor', 'Hospital', 'Date of Admission', 'Di
```

```
In [20]: from sklearn.preprocessing import LabelEncoder
```

```
In [21]: # Encode categorical variables
label_encoders = {}
for column in df.select_dtypes(include='object').columns:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column])
    label_encoders[column] = le
```

Classification: Predict Medical Condition

- Goal: Predict a patient's medical condition based on features like age, gender, blood type, test results, etc.
- ML Approach: Supervised Learning (Classification)

```
In [23]: from sklearn.model_selection import train_test_split
        from sklearn.metrics import classification_report, accuracy_score, confusion
        from imblearn.over_sampling import SMOTE
        from lightgbm import LGBMClassifier
```

```
In [24]: # Split features and target
        X = df.drop('Medical Condition', axis=1)
        y = df['Medical Condition']
```

```
In [25]: # Train/test split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```

```
In [26]: # Balance with SMOTE
        sm = SMOTE(random_state=42)
        X_res, y_res = sm.fit_resample(X_train, y_train)
```

```
In [27]: # Train LightGBM model
        model = LGBMClassifier(random_state=42)
        model.fit(X_res, y_res)
```

[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.003825 seconds.

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 359

[LightGBM] [Info] Number of data points in the train set: 44364, number of used features: 8

[LightGBM] [Info] Start training from score -1.791759

[LightGBM] [Info] Start training from score -1.791759

[LightGBM] [Info] Start training from score -1.791759

[LightGBM] [Info] Start training from score -1.791759

[LightGBM] [Info] Start training from score -1.791759

[LightGBM] [Info] Start training from score -1.791759

```
Out[27]: LGBMClassifier
         LGBMClassifier(random_state=42)
```

```
In [28]: # Predictions
        y_pred = model.predict(X_test)
```


```
In [29]: # Evaluation
        print("Accuracy:", accuracy_score(y_test, y_pred))
        print("Classification Report:\n", classification_report(y_test, y_pred))
        conf_matrix = confusion_matrix(y_test, y_pred)
```

Accuracy: 0.18782972530471165


Classification Report:

	precision	recall	f1-score	support
0	0.19	0.17	0.18	1867
1	0.19	0.19	0.19	1823
2	0.19	0.18	0.18	1830
3	0.19	0.20	0.19	1886
4	0.18	0.17	0.18	1757
5	0.19	0.22	0.21	1831
accuracy			0.19	10994
macro avg	0.19	0.19	0.19	10994
weighted avg	0.19	0.19	0.19	10994

```
In [30]: # Confusion matrix heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```
In [31]: # Feature importance
importance = pd.Series(model.feature_importances_, index=X.columns).sort_val
importance.plot(kind='barh', figsize=(10, 6), title="Feature Importances")
plt.gca().invert_yaxis()
plt.show()
```



```
In [32]: from sklearn.ensemble import RandomForestClassifier
```

```
In [33]: # Features and target
X = df.drop('Medical Condition', axis=1)
y = df['Medical Condition']

# Split into train/test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

# Train model
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# Predictions and evaluation
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred, output_dict=True)

# Feature importance
feature_importances = pd.Series(model.feature_importances_, index=X.columns)

accuracy, report, feature_importances
```

```

Out[33]: (0.27796980171002367,
{'0': {'precision': 0.2756892230576441,
      'recall': 0.2945902517407606,
      'f1-score': 0.28482651475919213,
      'support': 1867.0},
 '1': {'precision': 0.2679162072767365,
      'recall': 0.2665935271530444,
      'f1-score': 0.2672532306846302,
      'support': 1823.0},
 '2': {'precision': 0.28483491885842194,
      'recall': 0.27814207650273226,
      'f1-score': 0.2814487144042024,
      'support': 1830.0},
 '3': {'precision': 0.2838709677419355,
      'recall': 0.2799575821845175,
      'f1-score': 0.28190069407367857,
      'support': 1886.0},
 '4': {'precision': 0.2778702163061564,
      'recall': 0.28514513375071143,
      'f1-score': 0.28146067415730336,
      'support': 1757.0},
 '5': {'precision': 0.27780979827089336,
      'recall': 0.2632441288913162,
      'f1-score': 0.2703309029725182,
      'support': 1831.0},
 'accuracy': 0.27796980171002367,
 'macro avg': {'precision': 0.27799855525196465,
               'recall': 0.27794545003718035,
               'f1-score': 0.2778701218419208,
               'support': 10994.0},
 'weighted avg': {'precision': 0.2780279500336763,
                  'recall': 0.27796980171002367,
                  'f1-score': 0.2778962967812628,
                  'support': 10994.0}},
Billing Amount      0.342840
Age                 0.279717
Blood Type          0.108606
Insurance Provider  0.082574
Medication          0.082214
Test Results        0.039470
Admission Type      0.037921
Gender              0.026659
dtype: float64)

```

Admission Type Classification (Classification Problem)

- Predict the type of hospital admission using patient details.

Billing Amount Prediction (Regression Problem)

- Predict the billing amount based on patient and treatment features.

```

In [34]: from sklearn.preprocessing import LabelEncoder
         from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor

```

```

from sklearn.metrics import (
    accuracy_score, classification_report, confusion_matrix,
    mean_absolute_error, mean_squared_error, r2_score
)

```

```

In [35]: print("\n--- Admission Type Classification ---")

# Features and target
X_class = df.drop(columns=['Admission Type'])
y_class = df['Admission Type']
le_adm = LabelEncoder()
y_class_encoded = le_adm.fit_transform(y_class)

# Train-test split
X_train_cls, X_test_cls, y_train_cls, y_test_cls = train_test_split(X_class,

# Model training
cls_model = RandomForestClassifier(random_state=42)
cls_model.fit(X_train_cls, y_train_cls)

# Predictions
y_pred_cls = cls_model.predict(X_test_cls)

# Evaluation
print("Accuracy:", accuracy_score(y_test_cls, y_pred_cls))
print("Classification Report:\n", classification_report(y_test_cls, y_pred_c
print("Confusion Matrix:\n", confusion_matrix(y_test_cls, y_pred_cls))

# Feature Importance Plot
importances_cls = pd.Series(cls_model.feature_importances_, index=X_class.co
plt.figure(figsize=(10, 6))
sns.barplot(x=importances_cls[:10], y=importances_cls.index[:10])
plt.title("Top 10 Features - Admission Type Classification")
plt.xlabel("Importance Score")
plt.ylabel("Features")
plt.tight_layout()
plt.show()

```

--- Admission Type Classification ---

Accuracy: 0.4273239949063125

Classification Report:

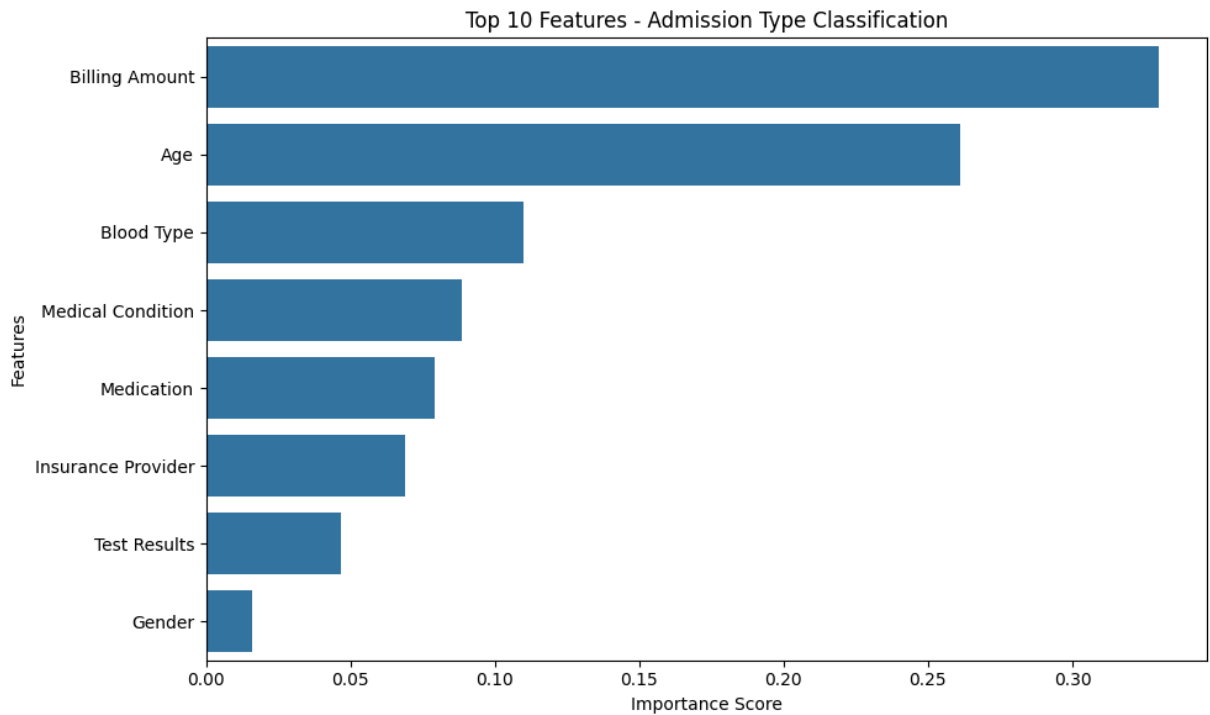
	precision	recall	f1-score	support
0	0.43	0.44	0.43	3665
1	0.43	0.41	0.42	3691
2	0.42	0.43	0.42	3638
accuracy			0.43	10994
macro avg	0.43	0.43	0.43	10994
weighted avg	0.43	0.43	0.43	10994

Confusion Matrix:

```

[[1615 1011 1039]
 [1063 1529 1099]
 [1098  986 1554]]

```



```
In [36]: print("\n--- Billing Amount Regression ---")

# Features and target
X_reg = df.drop(columns=['Billing Amount'])
y_reg = df['Billing Amount']

# Train-test split
X_train_reg, X_test_reg, y_train_reg, y_test_reg = train_test_split(X_reg, y

# Model training
reg_model = RandomForestRegressor(random_state=42)
reg_model.fit(X_train_reg, y_train_reg)

# Predictions
y_pred_reg = reg_model.predict(X_test_reg)

# Evaluation
print("MAE:", mean_absolute_error(y_test_reg, y_pred_reg))
print("RMSE:", mean_squared_error(y_test_reg, y_pred_reg, squared=False))
print("R^2 Score:", r2_score(y_test_reg, y_pred_reg))

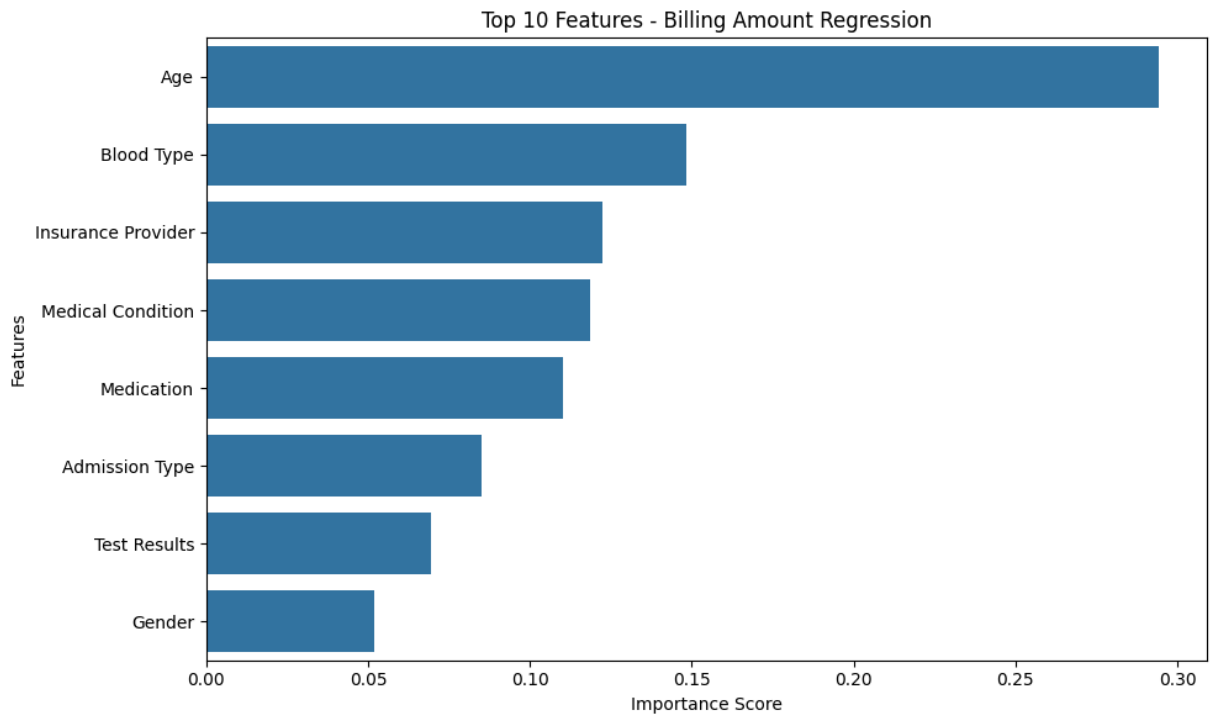
# Feature Importance Plot
importances_reg = pd.Series(reg_model.feature_importances_, index=X_reg.columns)
plt.figure(figsize=(10, 6))
sns.barplot(x=importances_reg[:10], y=importances_reg.index[:10])
plt.title("Top 10 Features - Billing Amount Regression")
plt.xlabel("Importance Score")
plt.ylabel("Features")
plt.tight_layout()
plt.show()
```

--- Billing Amount Regression ---

MAE: 12447.875469200017

RMSE: 14583.89163449932

R² Score: -0.04429357624767016



```
In [37]: from sklearn.model_selection import GridSearchCV
from xgboost import XGBClassifier
```

```
In [38]: # Features and target
X = df.drop(columns=['Admission Type'])
y = df['Admission Type']
le_target = LabelEncoder()
y_encoded = le_target.fit_transform(y)
```

```
In [39]: # Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=
```

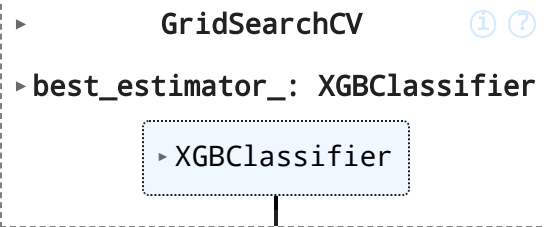
```
In [40]: # XGBoost model with hyperparameter tuning
xgb = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')

param_grid = {
    'n_estimators': [50, 100],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1.0]
}

grid_search = GridSearchCV(estimator=xgb, param_grid=param_grid, cv=3, scoring=
grid_search.fit(X_train, y_train)
```

Fitting 3 folds for each of 36 candidates, totalling 108 fits


```
Out[40]:
```



```

GridSearchCV
  best_estimator_: XGBClassifier
    XGBClassifier

```

```
In [41]: # Best model
best_model = grid_search.best_estimator_
```

```
In [42]: # Prediction
y_pred = best_model.predict(X_test)
```

```
In [43]: # Evaluation
print("Best Parameters:", grid_search.best_params_)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
Best Parameters: {'learning_rate': 0.2, 'max_depth': 7, 'n_estimators': 100,
'subsample': 0.8}
```

```
Accuracy: 0.36319810805894126
```

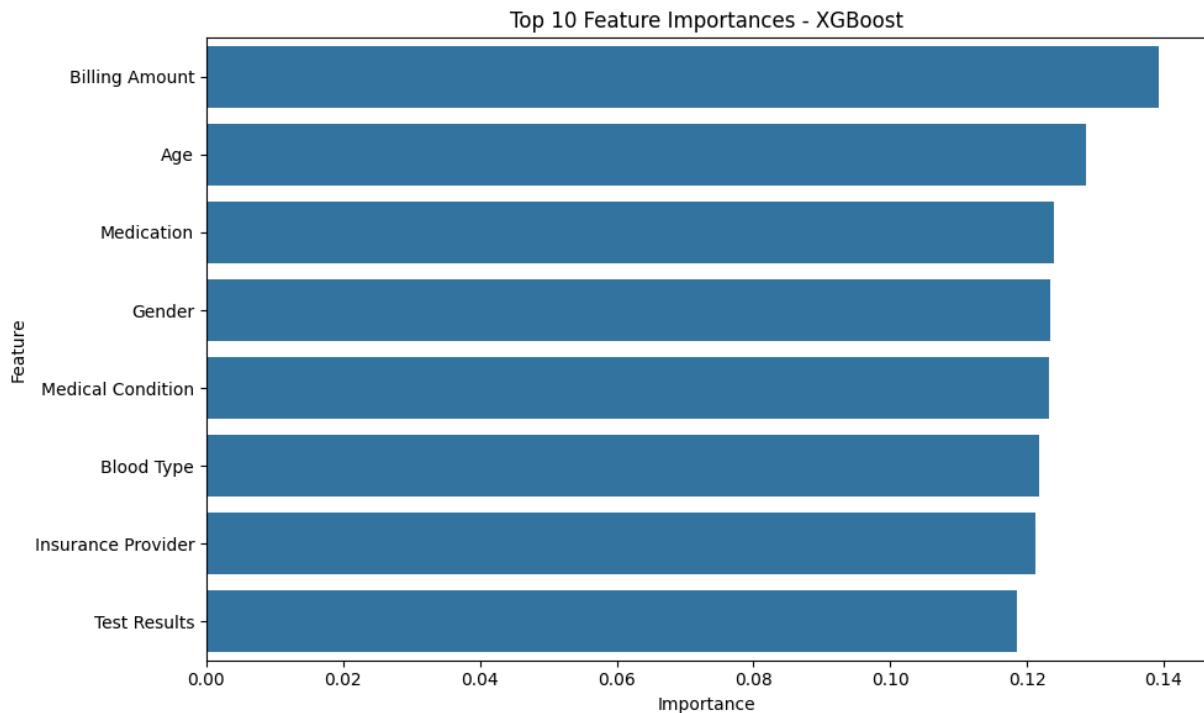
```
Classification Report:
```

	precision	recall	f1-score	support
0	0.36	0.38	0.37	3665
1	0.37	0.33	0.35	3691
2	0.37	0.39	0.38	3638
accuracy			0.36	10994
macro avg	0.36	0.36	0.36	10994
weighted avg	0.36	0.36	0.36	10994

```
Confusion Matrix:
```

```
[[1375 1083 1207]
 [1268 1216 1207]
 [1208 1028 1402]]
```

```
In [44]: # Feature importance plot
importances = pd.Series(best_model.feature_importances_, index=X.columns).so
plt.figure(figsize=(10, 6))
sns.barplot(x=importances[:10], y=importances.index[:10])
plt.title("Top 10 Feature Importances - XGBoost")
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



```
In [45]: from sklearn.model_selection import train_test_split, StratifiedKFold
        from xgboost import plot_importance
```

```
In [46]: # Stratified K-Fold for better validation
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Expanded parameter grid
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.05, 0.1],
    'subsample': [0.7, 0.9, 1.0],
    'colsample_bytree': [0.7, 0.9, 1.0],
    'gamma': [0, 1, 5]
}

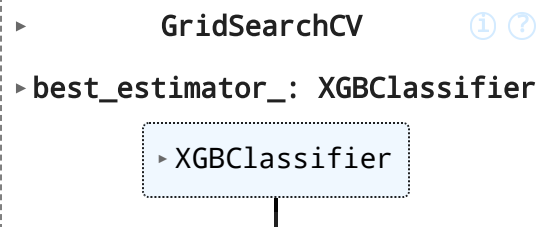
xgb = XGBClassifier(
    objective='multi:softprob',
    use_label_encoder=False,
    eval_metric='mlogloss',
    random_state=42
)

grid_search = GridSearchCV(
    estimator=xgb,
    param_grid=param_grid,
    cv=cv,
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)
```

```
In [47]: # Fit with early stopping
grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 486 candidates, totalling 2430 fits

```
Out[47]:
```



```
In [48]: # Best model
best_model = grid_search.best_estimator_
print("Best Parameters:", grid_search.best_params_)
```

Best Parameters: {'colsample_bytree': 1.0, 'gamma': 0, 'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 200, 'subsample': 0.7}

```
In [49]: # Accuracy
y_pred = best_model.predict(X_test)
print("Test Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

Test Accuracy: 0.3661997453156267

Classification Report:

	precision	recall	f1-score	support
0	0.36	0.38	0.37	3665
1	0.38	0.34	0.36	3691
2	0.37	0.37	0.37	3638
accuracy			0.37	10994
macro avg	0.37	0.37	0.37	10994
weighted avg	0.37	0.37	0.37	10994

Confusion Matrix:
[[1410 1054 1201]
[1282 1257 1152]
[1260 1019 1359]]

```
In [60]: # More Advance Working Sonn..!
```

```
In [ ]: # Drop irrelevant columns
# df = df.drop(columns=['Name', 'Doctor', 'Hospital', 'Date of Admission', '])
```

```
In [61]: # Apply map() to encode categorical variables
mappings = {}
for col in df.select_dtypes(include='object').columns:
    unique_vals = df[col].unique()
    mapping = {val: idx for idx, val in enumerate(unique_vals)}
    df[col] = df[col].map(mapping)
    mappings[col] = mapping
```

```

# Features & target
X = df.drop(columns=['Admission Type', 'Insurance Provider', 'Billing Amount'])
y = df['Admission Type']
target_mapping = {val: idx for idx, val in enumerate(y.unique())}
y = y.map(target_mapping)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_s

```

```

In [62]: # XGBoost model with hyperparameter tuning
xgb = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')

param_grid = {
    'n_estimators': [50, 100],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1.0]
}

grid_search = GridSearchCV(estimator=xgb, param_grid=param_grid, cv=3, scoring=
grid_search.fit(X_train, y_train)

```

Fitting 3 folds for each of 36 candidates, totalling 108 fits

```

Out[62]:
GridSearchCV
└─ best_estimator_: XGBClassifier
    └─ XGBClassifier

```

```

In [63]: # Best model
best_model = grid_search.best_estimator_

```

```

In [64]: # Prediction
y_pred = best_model.predict(X_test)

```

```

In [65]: # Evaluation
print("Best Parameters:", grid_search.best_params_)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

Best Parameters: {'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 100, 'subsample': 0.8}

Accuracy: 0.34027651446243407

Classification Report:

	precision	recall	f1-score	support
0	0.34	0.34	0.34	3678
1	0.34	0.30	0.32	3621
2	0.34	0.38	0.36	3695
accuracy			0.34	10994
macro avg	0.34	0.34	0.34	10994
weighted avg	0.34	0.34	0.34	10994

Confusion Matrix:

```
[[1253 1086 1339]
 [1188 1085 1348]
 [1258 1034 1403]]
```

```
In [72]: df['Admission Type'].value_counts()
```

```
Out[72]: Admission Type
Elective      18655
Urgent        18576
Emergency     18269
Name: count, dtype: int64
```

```
In [77]: y_test.value_counts()
```

```
Out[77]: Admission Type
2      3695
0      3678
1      3621
Name: count, dtype: int64
```

```
In [16]: # Admission Type Classification (XGBoost + Map Encoding)
```

```
In [57]: from sklearn.model_selection import train_test_split, StratifiedKFold, GridS
from sklearn.metrics import accuracy_score, classification_report, confusion
from xgboost import XGBClassifier, XGBRegressor, plot_importance
from imblearn.over_sampling import SMOTE
```

```
In [7]: # Map Gender
df['Gender'] = df['Gender'].map({'Male': 1, 'Female': 0})

# Map Blood Type
df['Blood Type'] = df['Blood Type'].map({
    'A-': 0, 'A+': 1, 'B+': 2, 'AB+': 3, 'AB-': 4, 'B-': 5, 'O+': 6, 'O-': 7
})

# Map Medical Condition
df['Medical Condition'] = df['Medical Condition'].map({
    'Arthritis': 0, 'Diabetes': 1, 'Hypertension': 2,
    'Obesity': 3, 'Cancer': 4, 'Asthma': 5
})
```

```

# Map Insurance Provider
df['Insurance Provider'] = df['Insurance Provider'].map({
    'Cigna': 0, 'Medicare': 1, 'UnitedHealthcare': 2,
    'Blue Cross': 3, 'Aetna': 4
})

# Map Medication
df['Medication'] = df['Medication'].map({
    'Lipitor': 0, 'Ibuprofen': 1, 'Aspirin': 2,
    'Paracetamol': 3, 'Penicillin': 4
})

# Map Test Results
df['Test Results'] = df['Test Results'].map({
    'Abnormal': 0, 'Normal': 1, 'Inconclusive': 2
})

# Map target variable Admission Type
df['Admission Type'] = df['Admission Type'].map({
    'Elective': 0, 'Urgent': 1, 'Emergency': 2
})

```

```

In [38]: # Select relevant features and target
features = ['Age', 'Gender', 'Blood Type', 'Medical Condition', 'Insurance P
target = 'Admission Type'

```

```

X = df[features]
y = df[target]

```

```

In [39]: # Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_s

```

```

In [40]: # Apply SMOTE to balance classes
smote = SMOTE(random_state=42)
X_train_res, y_train_res = smote.fit_resample(X_train, y_train)

```

```

In [41]: # Define XGBoost classifier
xgb = XGBClassifier(
    objective='multi:softprob',
    use_label_encoder=False,
    eval_metric='mlogloss',
    random_state=42
)

# Hyperparameter grid
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5],
    'learning_rate': [0.05, 0.1],
    'subsample': [0.8, 1.0],
    'colsample_bytree': [0.8, 1.0],
    'gamma': [0, 1]
}

# Grid search with cross-validation

```

```

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
grid_search = GridSearchCV(
    estimator=xgb,
    param_grid=param_grid,
    cv=cv,
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)

```

```

In [42]: # Train the model
grid_search.fit(X_train_res, y_train_res)
best_model = grid_search.best_estimator_

```

Fitting 5 folds for each of 64 candidates, totalling 320 fits

```

In [43]: # Predict and evaluate
y_pred = best_model.predict(X_test)

```

```

In [44]: print("\nBest Parameters:", grid_search.best_params_)
print("Test Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

Best Parameters: {'colsample_bytree': 0.8, 'gamma': 0, 'learning_rate': 0.1, 'max_depth': 5, 'n_estimators': 200, 'subsample': 0.8}

Test Accuracy: 0.34300527560487537

Classification Report:

	precision	recall	f1-score	support
0	0.35	0.34	0.35	3695
1	0.34	0.35	0.34	3678
2	0.34	0.34	0.34	3621
accuracy			0.34	10994
macro avg	0.34	0.34	0.34	10994
weighted avg	0.34	0.34	0.34	10994

Confusion Matrix:

```

[[1261 1282 1152]
 [1185 1277 1216]
 [1156 1232 1233]]

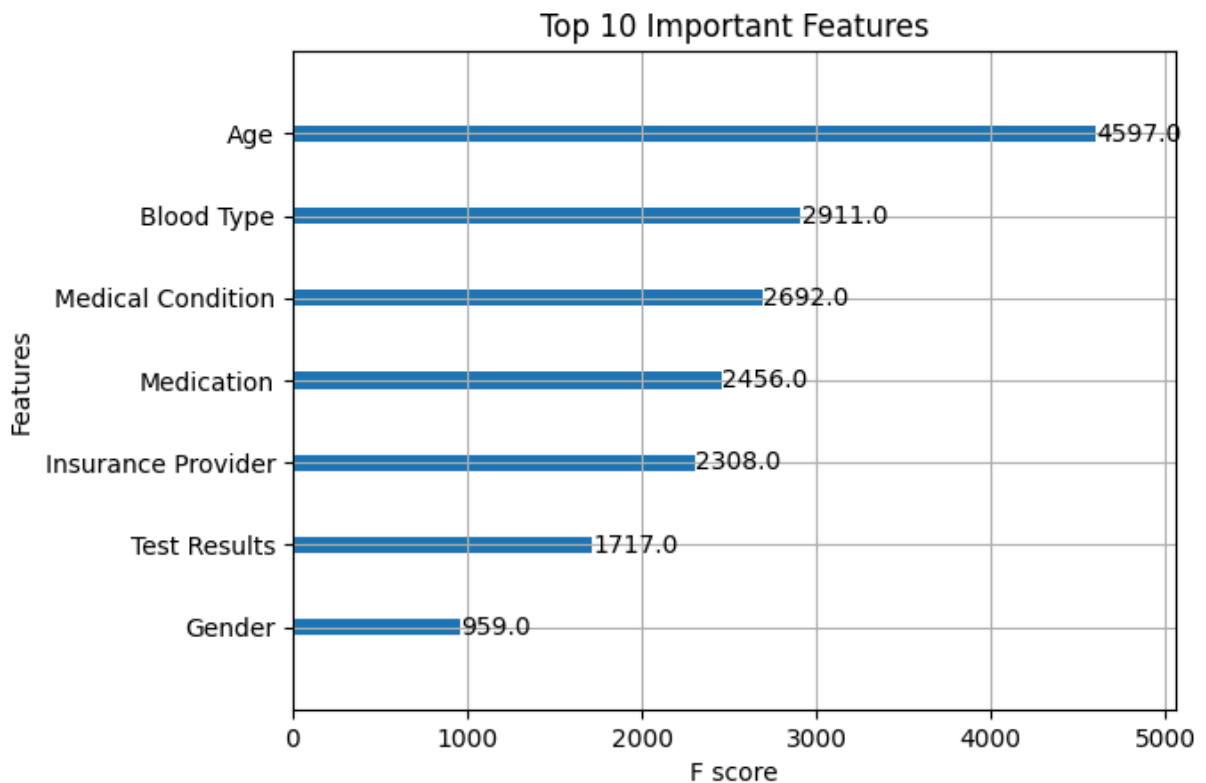
```

```

In [45]: # Plot Feature Importance
plt.figure(figsize=(10, 6))
plot_importance(best_model, max_num_features=10)
plt.title("Top 10 Important Features")
plt.show()

```

<Figure size 1000x600 with 0 Axes>



```
In [47]: import numpy as np
# Assuming 'best_model' is already loaded or trained

print('--- Enter the Petition Details ---')

# Taking inputs for features
age = int(input('Enter the Age of the Patient: '))
gender = int(input('Enter the Gender (0: Female, 1: Male): '))
blood_type = int(input('Enter the Blood Type (0: A, 1: B, 2: AB, 3: O): '))
medical_condition = float(input('Enter the Medical Condition Score (e.g., 0.5): '))
medication = int(input('Enter Medication Code (e.g., 0 or 1): '))
test_results = int(input('Enter the Test Results Score (e.g., 0 or 1): '))
insurance_provider = int(input('Enter the Insurance Provider Code: '))
# billing_amount = float(input('Enter the Billing Amount (in USD): '))

# Preparing input for the model
input_point = np.array([[age, gender, blood_type, medical_condition, medication,
                          test_results, insurance_provider]]) # billing_amount

# Making prediction
prediction = best_model.predict(input_point)

# Mapping prediction to human-readable label
label_mapping = {0: 'Elective', 1: 'Urgent', 2: 'Emergency'}
predicted_label = label_mapping.get(prediction[0], "Unknown")

# Printing result
print(f'\nPrediction Result: {predicted_label} Admission')
```



```
--- Enter the Petition Details ---
Enter the Age of the Patient: 30
Enter the Gender (0: Female, 1: Male): 1
Enter the Blood Type (0: A, 1: B, 2: AB, 3: O): 5
Enter the Medical Condition Score (e.g., 0.0 to 1.0): 3
Enter Medication Code (e.g., 0 or 1): 3
Enter the Test Results Score (e.g., 0 or 1): 1
Enter the Insurance Provider Code: 2
```

Prediction Result: Emergency Admission

```
In [48]: # Features and Targets
features = ['Age', 'Gender', 'Blood Type', 'Medical Condition', 'Insurance P
X = df[features]
y_class = df['Admission Type']
y_reg = df['Billing Amount']

In [49]: # Split for classification
X_train_class, X_test_class, y_train_class, y_test_class = train_test_split(

In [50]: # Apply SMOTE to balance Admission Type
smote = SMOTE(random_state=42)
X_train_class_res, y_train_class_res = smote.fit_resample(X_train_class, y_t

In [52]: # XGBoost Classifier
xgb_clf = XGBClassifier(
    objective='multi:softmax',
    num_class=3,
    eval_metric='mlogloss',
    use_label_encoder=False,
    random_state=42
)
xgb_clf.fit(X_train_class_res, y_train_class_res)
y_pred_class = xgb_clf.predict(X_test_class)

In [53]: # Evaluate classifier
print("🔥 Admission Type Classification")
print("Accuracy:", accuracy_score(y_test_class, y_pred_class))
print("Classification Report:\n", classification_report(y_test_class, y_pred
print("Confusion Matrix:\n", confusion_matrix(y_test_class, y_pred_class))
```

🔴 Admission Type Classification

Accuracy: 0.3369110423867564

Classification Report:

	precision	recall	f1-score	support
0	0.34	0.34	0.34	3695
1	0.34	0.33	0.34	3678
2	0.33	0.34	0.34	3621
accuracy			0.34	10994
macro avg	0.34	0.34	0.34	10994
weighted avg	0.34	0.34	0.34	10994

Confusion Matrix:

```
[[1242 1213 1240]
 [1229 1232 1217]
 [1225 1166 1230]]
```

```
In [54]: # Split for regression
X_train_reg, X_test_reg, y_train_reg, y_test_reg = train_test_split(X, y_reg
```

```
In [55]: # XGBoost Regressor
xgb_reg = XGBRegressor(
    objective='reg:squarederror',
    n_estimators=100,
    learning_rate=0.1,
    max_depth=5,
    random_state=42
)
xgb_reg.fit(X_train_reg, y_train_reg)
y_pred_reg = xgb_reg.predict(X_test_reg)
```

```
In [58]: # Evaluate regression
print("\n🔴 Billing Amount Prediction")
print("MAE:", mean_absolute_error(y_test_reg, y_pred_reg))
print("RMSE:", np.sqrt(mean_squared_error(y_test_reg, y_pred_reg)))
print("R² Score:", r2_score(y_test_reg, y_pred_reg))
```

🔴 Billing Amount Prediction

MAE: 12408.12414265777

RMSE: 14332.809098201316


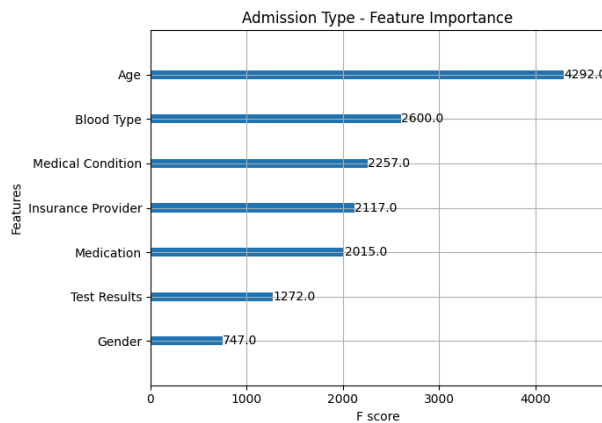
R² Score: -0.00864509790921808

```
In [59]: # Plot feature importances for both models
plt.figure(figsize=(14, 5))

plt.subplot(1, 2, 1)
plot_importance(xgb_clf, ax=plt.gca(), title='Admission Type - Feature Importances')

plt.subplot(1, 2, 2)
plot_importance(xgb_reg, ax=plt.gca(), title='Billing Amount - Feature Importances')

plt.tight_layout()
plt.show()
```



```
In [ ]: # Classification + Regression + Input Prediction
```

```
In [62]: # Features and targets
features = ['Age', 'Gender', 'Medical Condition', 'Insurance Provider', 'Med
X = df[features]
y_class = df['Admission Type']
y_reg = df['Billing Amount']

# SMOTE for classification
smote = SMOTE(random_state=42)
X_smote, y_class_smote = smote.fit_resample(X, y_class)

# Split data
X_train_class, X_test_class, y_train_class, y_test_class = train_test_split(
X_train_reg, X_test_reg, y_train_reg, y_test_reg = train_test_split(X, y_reg

# XGBoost Classifier with hyperparameter tuning
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.05, 0.1, 0.2]
}
grid_search = GridSearchCV(XGBClassifier(random_state=42, use_label_encoder=
grid_search.fit(X_train_class, y_train_class)
best_model_class = grid_search.best_estimator_

# Evaluate classifier
y_pred_class = best_model_class.predict(X_test_class)
print("\n--- Classification Report (Admission Type) ---")
print(classification_report(y_test_class, y_pred_class))

# XGBoost Regressor for billing amount
best_model_reg = XGBRegressor(n_estimators=100, max_depth=5, learning_rate=0
best_model_reg.fit(X_train_reg, y_train_reg)

# Evaluate regressor
y_pred_reg = best_model_reg.predict(X_test_reg)
rmse = np.sqrt(mean_squared_error(y_test_reg, y_pred_reg))
print(f"\n--- Regression RMSE (Billing Amount): ${rmse:.2f}")
```

```

--- Classification Report (Admission Type) ---
              precision    recall  f1-score   support

     0           0.34         0.34         0.34         3703
     1           0.33         0.35         0.34         3626
     2           0.35         0.34         0.34         3755

 accuracy              0.34         0.34         0.34         11084
 macro avg           0.34         0.34         0.34         11084
 weighted avg        0.34         0.34         0.34         11084

```

```

--- Regression RMSE (Billing Amount): $14321.94

```

```

In [63]: # Prediction Code (User Input + Output)
# 🧠 Prediction section

print('--- Enter Patient Details ---')

age = int(input('Enter the Age of the Patient: '))
gender = int(input('Enter the Gender (0: Female, 1: Male): '))
medical_condition = int(input('Enter Medical Condition Code (0: Arthritis, 1: Diabetes, 2: Hypertension, 3: Obesity, 4: Cancer, 5: Asthma): '))
insurance_provider = int(input('Enter Insurance Provider Code (0: Cigna, 1: Medicare, 2: UnitedHealthcare, 3: Blue Cross, 4: Aetna): '))
medication = int(input('Enter Medication Code (0: Lipitor, 1: Ibuprofen, 2: Aspirin, 3: Paracetamol, 4: Penicillin): '))
test_results = int(input('Enter Test Results Code (0: Abnormal, 1: Normal, 2: Inconclusive): '))

# Prepare input for prediction
input_point = np.array([age, gender, medical_condition, insurance_provider, medication, test_results])

# Predictions
admission_prediction = best_model_class.predict(input_point)
billing_prediction = best_model_reg.predict(input_point)

# Mapping Admission Type
label_mapping = {0: 'Elective', 1: 'Urgent', 2: 'Emergency'}
admission_type = label_mapping.get(admission_prediction[0], "Unknown")

# Final Output
print('\n🔍 Predicted Results:')
print(f'👨‍⚕️ Admission Type: {admission_type}')
print(f'💰 Estimated Billing Amount: ${billing_prediction[0]:.2f}')

```

```

--- Enter Patient Details ---
Enter the Age of the Patient: 30
Enter the Gender (0: Female, 1: Male): 1
Enter Medical Condition Code (0: Arthritis, 1: Diabetes, 2: Hypertension, 3: Obesity, 4: Cancer, 5: Asthma): 4
Enter Insurance Provider Code (0: Cigna, 1: Medicare, 2: UnitedHealthcare, 3: Blue Cross, 4: Aetna): 3
Enter Medication Code (0: Lipitor, 1: Ibuprofen, 2: Aspirin, 3: Paracetamol, 4: Penicillin): 3
Enter Test Results Code (0: Abnormal, 1: Normal, 2: Inconclusive): 1

```

```

🔍 Predicted Results:
👨‍⚕️ Admission Type: Elective
💰 Estimated Billing Amount: $26140.86

```

```
In [ ]: # Blending Machine Learning Method Apply
```

```
In [27]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import classification_report, mean_squared_error, accur
from sklearn.ensemble import GradientBoostingClassifier, RandomForestClassif
from sklearn.ensemble import GradientBoostingRegressor, RandomForestRegresso
from imblearn.over_sampling import SMOTE
from xgboost import XGBClassifier, XGBRegressor
```


```
In [15]: # Mapping categorical columns
df['Gender'] = df['Gender'].map({'Male': 1, 'Female': 0})
df['Blood Type'] = df['Blood Type'].map({'A-': 0, 'A+': 1, 'B+': 2, 'AB+': 3
df['Medical Condition'] = df['Medical Condition'].map({'Arthritis': 0, 'Diab
df['Insurance Provider'] = df['Insurance Provider'].map({'Cigna': 0, 'Medica
df['Medication'] = df['Medication'].map({'Lipitor': 0, 'Ibuprofen': 1, 'Aspi
df['Test Results'] = df['Test Results'].map({'Abnormal': 0, 'Normal': 1, 'In
df['Admission Type'] = df['Admission Type'].map({'Elective': 0, 'Urgent': 1,
```

```
In [16]: # Features and targets
features = ['Age', 'Gender', 'Blood Type', 'Medical Condition', 'Insurance P
X = df[features]
y_class = df['Admission Type']
y_reg = df['Billing Amount']
```

```
In [17]: # Apply SMOTE
smote = SMOTE(random_state=42)
X_smote, y_class_smote = smote.fit_resample(X, y_class)
```

```
In [18]: # Train/test split
X_train_cls, X_test_cls, y_train_cls, y_test_cls = train_test_split(X_smote,
X_train_reg, X_test_reg, y_train_reg, y_test_reg = train_test_split(X, y_reg
```

```
In [19]: # Blending for Classification
clf1 = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss', random
clf2 = GradientBoostingClassifier(random_state=42)
clf3 = RandomForestClassifier(random_state=42)
voting_clf = VotingClassifier(estimators=[('xgb', clf1), ('gb', clf2), ('rf'
voting_clf.fit(X_train_cls, y_train_cls)
```



```
In [20]: # Blending for Regression
reg1 = XGBRegressor(random_state=42)
reg2 = GradientBoostingRegressor(random_state=42)
reg3 = RandomForestRegressor(random_state=42)
```

```
In [21]: # Fit individual regressors
reg1.fit(X_train_reg, y_train_reg)
reg2.fit(X_train_reg, y_train_reg)
reg3.fit(X_train_reg, y_train_reg)
```

```
Out[21]: RandomForestRegressor
RandomForestRegressor(random_state=42)
```

```
In [22]: # Evaluate classification
y_pred_cls = voting_clf.predict(X_test_cls)
cls_report = classification_report(y_test_cls, y_pred_cls, output_dict=True)
```

```
In [23]: # Evaluate regression (average ensemble)
y_pred_reg1 = reg1.predict(X_test_reg)
y_pred_reg2 = reg2.predict(X_test_reg)
y_pred_reg3 = reg3.predict(X_test_reg)
y_pred_reg_avg = (y_pred_reg1 + y_pred_reg2 + y_pred_reg3) / 3
reg_rmse = np.sqrt(mean_squared_error(y_test_reg, y_pred_reg_avg))
```

```
In [24]: cls_report, reg_rmse
```

```
Out[24]: ({'0': {'precision': 0.38834688346883467,
  'recall': 0.38698352687010534,
  'f1-score': 0.38766400649262817,
  'support': 3703.0},
  '1': {'precision': 0.3824317086234601,
  'recall': 0.3938223938223938,
  'f1-score': 0.38804347826086955,
  'support': 3626.0},
  '2': {'precision': 0.3983606557377049,
  'recall': 0.3882822902796272,
  'f1-score': 0.39325691166554283,
  'support': 3755.0},
  'accuracy': 0.38966077228437385,
  'macro avg': {'precision': 0.3897130826099999,
  'recall': 0.38969607032404213,
  'f1-score': 0.38965479880634685,
  'support': 11084.0},
  'weighted avg': {'precision': 0.3898042355872287,
  'recall': 0.38966077228437385,
  'f1-score': 0.3896828916925503,
  'support': 11084.0}},
 14367.02594139383)
```

```
In [28]: # Train/test split
X_train_cls, X_test_cls, y_train_cls, y_test_cls = train_test_split(X_smote,
X_train_reg, X_test_reg, y_train_reg, y_test_reg = train_test_split(X, y_reg

# Classifiers
clf1 = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss', random
clf2 = GradientBoostingClassifier(random_state=42)
clf3 = RandomForestClassifier(random_state=42)
```

```

voting_clf = VotingClassifier(estimators=[('xgb', clf1), ('gb', clf2), ('rf',
# Train classifiers
clf1.fit(X_train_cls, y_train_cls)
clf2.fit(X_train_cls, y_train_cls)
clf3.fit(X_train_cls, y_train_cls)
voting_clf.fit(X_train_cls, y_train_cls)

# Predict and evaluate classifiers
models_cls = {'XGBoost': clf1, 'GradientBoosting': clf2, 'RandomForest': clf3}
acc_scores = {}

print("\n--- Classification Accuracy Scores ---")
for name, model in models_cls.items():
    preds = model.predict(X_test_cls)
    acc = accuracy_score(y_test_cls, preds)
    acc_scores[name] = acc
    print(f"{name}: Accuracy = {acc:.4f}")

best_cls_model = max(acc_scores, key=acc_scores.get)
print(f"\nBest Classification Model: {best_cls_model} with Accuracy = {acc_s

# Regressors
reg1 = XGBRegressor(random_state=42)
reg2 = GradientBoostingRegressor(random_state=42)
reg3 = RandomForestRegressor(random_state=42)

# Train regressors
reg1.fit(X_train_reg, y_train_reg)
reg2.fit(X_train_reg, y_train_reg)
reg3.fit(X_train_reg, y_train_reg)

# Predict and evaluate regressors
models_reg = {'XGBoost': reg1, 'GradientBoosting': reg2, 'RandomForest': reg3}
rmse_scores = {}

print("\n--- Regression RMSE Scores ---")
for name, model in models_reg.items():
    preds = model.predict(X_test_reg)
    rmse = np.sqrt(mean_squared_error(y_test_reg, preds))
    rmse_scores[name] = rmse
    print(f"{name}: RMSE = {rmse:.2f}")

# Average ensemble prediction
y_pred_reg_avg = (reg1.predict(X_test_reg) + reg2.predict(X_test_reg) + reg3
rmse_avg = np.sqrt(mean_squared_error(y_test_reg, y_pred_reg_avg))
rmse_scores['AverageEnsemble'] = rmse_avg
print(f"Average Ensemble RMSE = {rmse_avg:.2f}")

best_reg_model = min(rmse_scores, key=rmse_scores.get)
print(f"\nBest Regression Model: {best_reg_model} with RMSE = {rmse_scores[b

# Optional: Plot scores
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.bar(acc_scores.keys(), acc_scores.values(), color='skyblue')

```

```
plt.title("Classification Accuracy")
plt.ylabel("Accuracy")
plt.xticks(rotation=45)

plt.subplot(1, 2, 2)
plt.bar(rmse_scores.keys(), rmse_scores.values(), color='lightgreen')
plt.title("Regression RMSE")
plt.ylabel("RMSE")
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

--- Classification Accuracy Scores ---

XGBoost: Accuracy = 0.3436

GradientBoosting: Accuracy = 0.3421

RandomForest: Accuracy = 0.3990

VotingEnsemble: Accuracy = 0.3897

Best Classification Model: RandomForest with Accuracy = 0.3990

--- Regression RMSE Scores ---

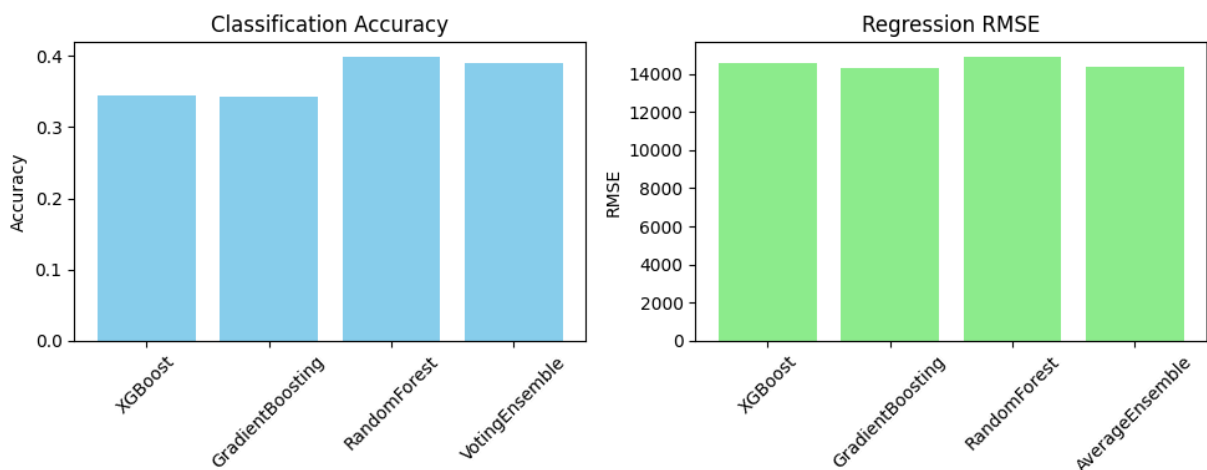
XGBoost: RMSE = 14554.68

GradientBoosting: RMSE = 14286.19

RandomForest: RMSE = 14908.73

Average Ensemble RMSE = 14367.03

Best Regression Model: GradientBoosting with RMSE = 14286.19



```
In [33]: # Prediction Code (User Input + Output)
# 🤖 Prediction section

print('--- Enter Patient Details ---')

age = int(input('Enter the Age of the Patient: '))
gender = int(input('Enter the Gender (0: Female, 1: Male): '))
blood_type = int(input('Enter Blood Type Code (0: A-, 1: A+, 2: B+, 3: AB+, 4: O-): '))
medical_condition = int(input('Enter Medical Condition Code (0: Arthritis, 1: Diabetes, 2: Hypertension, 3: Heart Disease, 4: Kidney Disease, 5: Liver Disease, 6: Multiple Sclerosis, 7: Parkinson's Disease, 8: Rheumatoid Arthritis, 9: Stroke): '))
insurance_provider = int(input('Enter Insurance Provider Code (0: Cigna, 1: Aetna, 2: UnitedHealthcare, 3: Anthem, 4: Humana, 5: Kaiser Permanente, 6: Blue Cross Blue Shield, 7: Centene, 8: Wellpoint, 9: Cofinity): '))
medication = int(input('Enter Medication Code (0: Lipitor, 1: Ibuprofen, 2: Metformin, 3: Lisinopril, 4: Zolpidem, 5: Sertraline, 6: Fentanyl, 7: Insulin, 8: Warfarin, 9: Clozapine): '))
test_results = int(input('Enter Test Results Code (0: Abnormal, 1: Normal, 2: Borderline): '))
```



```

# Prepare input for prediction
input_point = np.array([age, gender, blood_type, medical_condition, insurance_provider])

# Predictions
admission_prediction = clf3.predict(input_point)
billing_prediction = reg2.predict(input_point)

# Mapping Admission Type
label_mapping = {0: 'Elective', 1: 'Urgent', 2: 'Emergency'}
admission_type = label_mapping.get(admission_prediction[0], "Unknown")

# Final Output
print('\n🔍 Predicted Results:')
print(f'🏥 Admission Type: {admission_type}')
print(f'💰 Estimated Billing Amount: ₹{billing_prediction[0]:.2f}')

```

--- Enter Patient Details ---

Enter the Age of the Patient: 30

Enter the Gender (0: Female, 1: Male): 1

Enter Blood Type Code (0: A-, 1: A+, 2: B+, 3: AB+, 4: AB-, 5: B-, 6: O+, 7: O-): 5

Enter Medical Condition Code (0: Arthritis, 1: Diabetes, 2: Hypertension, 3: Obesity, 4: Cancer, 5: Asthma): 4

Enter Insurance Provider Code (0: Cigna, 1: Medicare, 2: UnitedHealthcare, 3: Blue Cross, 4: Aetna): 3

Enter Medication Code (0: Lipitor, 1: Ibuprofen, 2: Aspirin, 3: Paracetamol, 4: Penicillin): 3

Enter Test Results Code (0: Abnormal, 1: Normal, 2: Inconclusive): 1

🔍 Predicted Results:

🏥 Admission Type: Urgent

💰 Estimated Billing Amount: ₹26061.43

In []: *# More Advance Working Sonn..!*
Notebook Project By : PRASAD JADHAV (ML-ENG)
LinkedIn: [linkedin.com/in/prasadmjadhav2](https://www.linkedin.com/in/prasadmjadhav2) | Github: github.com/prasadmjadhav2

Thank You!